

1. Intro to text cleaning
2. REGEX exercise

Dr. Noa Cohen

Regular Expressions

- A **regular expression** is a kind of **pattern** that can be applied to text (Strings, in Python)
- A regular expression either **matches** the text (or part of the text), or it fails to match it
 - **Which part** of the text matches?
 - **Which parts** of the regular expression match **which parts** of the matched text?
 - Can we **perform substitutions** on the text?
- Regular expressions are an extremely useful tool for manipulating text

Machine Learning and Texts

- Natural language processing (NLP) application examples:
 - Autocorrect – correct to a valid word, correct to a valid word in the context of the sentence
 - Rephrasing of a sentence (see wordtune.com)
 - Recognizing harmful conversations
 - Emails classification (spam \ not spam)
 - Translation between languages

Language Models

- Language models (LM) are an essential element in Natural Language Processing (NLP)
- Many popular NLP applications are using language models
 - Google Assistant, Siri, Amazon's Alexa
 - Text summarization
 - Predicting which word should come next (Google's autofill)
 - Speech recognition
 - Machine translation
 - Optical Character Recognition (OCR)
 - Many more

Language Models

- A language model predicts the probability of a sentence (a sequence of words)
- For example, let us consider Google Translate. In Machine Translation, we convert a set of words from one language to another
 - Often, there are many potential translations
 - Using the probability of each translation to be valid, we can prioritize

$p(\text{"the class is interesting"}) > p(\text{"interesting the is class"})$

$p(\text{"the class is interesting"}) >$

$p(\text{"interesting the class is"}) >$

$p(\text{"interesting the is class"})$

"the class is interesting"

Q All Images

About 1,650,000 results

"interesting the class is"

Q All Images

About 280,000 results (0

"interesting the is class"

Did you mean: "interesting *this* class"

No results found for "interesting the is class".

Language Models

- In the presented example, by knowing that the probability of the first sentence is higher than the rest, we arrive at the right translation
- This ability to model the rules of a language as a **probability** gives great power for NLP related tasks

Text Cleaning

- When working with texts, we can not build language models or train machine\deep learning models on the raw data
- I.e., we should start be pre-processing of cleaning the text first
 - Ideas for “cleaning the text”?
 - Splitting it into words, handling punctuation, handling case, handling numbers, and more
- There is a set of text preparation methods that we should consider using; the choice of methods depends the specific natural language processing (NLP) task

Text Cleaning - motivation

- Assuming we want to count the number of appearances of the sentence “the class is interesting” in the text
- It could appear as different variations:
 - As a title: “The Class is Interesting” or “THE CLASS IS INTERESTING”
 - With extra spaces: “the class is interesting”
 - With a newline: “the class is\ninteresting”
 - With html tags (if we downloaded a web page):
“the class is interesting”
 - With punctuation: “the class, is interesting!”
- → Text cleaning is required before we perform NLP analysis

Text Cleaning

- In python we can fetch all punctuations using

```
import string
```

```
string.punctuation
```

Exercise - REGEX

- Given a text file (see “TheLittlePrinceRegexExercise.txt”), find all:
 - Appearances of the word I
 - Report how many are found
 - Words starting with a capital letter (e.g., *The*, *Little*, *PRINCE*)
 - Report how many are found
 - Words separated by dashes (e.g., *air-planes*, *grown-up*)
 - Report how many are found in general
 - Keep them in a dictionary with number of appearances per expression (for example: {'air-planes': 1, 'grown-ups': 6, ...})

Exercise - REGEX

- Numbers (e.g., *1943*)
 - Report how many are found
 - Keep a list of these numbers sorted by order of appearances
- Adjacent duplicated words that appear twice (e.g., *and and*)
 - Report how many are found
 - Keep them in a dictionary with number of appearances
 - Print them to a file named “duplications.txt”
- Quotations, i.e., sentences surrounded with “” or “”
 - Report how many are found
 - Print them to a file named “quotations.txt”
- Report your answers to the shared file:
“REGEX_exercise_shared”

Exercise – Text Cleaning

- Write your own `clean_text()` function and apply it on the given data (see “TheLittlePrince.txt” in the course’s site)
- Generate the file “TheLittlePrinceCleaned.txt”:
 - Only lower-case letters
 - No punctuations
 - Handle new lines
 - Adjacent duplicated words (could be a mistake in the original text) – replace by one